

# Machine Learning & Data Mining

## CS/CNS/EE 155

Lecture 15:  
Hidden Markov Models

# Sequence Prediction (POS Tagging)

- $x = \text{"Fish Sleep"}$
- $y = (\text{N}, \text{V})$
  
- $x = \text{"The Dog Ate My Homework"}$
- $y = (\text{D}, \text{N}, \text{V}, \text{D}, \text{N})$
  
- $x = \text{"The Fox Jumped Over The Fence"}$
- $y = (\text{D}, \text{N}, \text{V}, \text{P}, \text{D}, \text{N})$

# Challenges

- Multivariable Output
  - Make multiple predictions simultaneously
- Variable Length Input/Output
  - Sentence lengths not fixed

# Multivariate Outputs

- $x = \text{"Fish Sleep"}$
- $y = (N, V)$
- Multiclass prediction:

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

Replicate Weights:

$$w = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_K \end{bmatrix} \quad b = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_K \end{bmatrix}$$

Score All Classes:

$$f(x | w, b) = \begin{bmatrix} w_1^T x - b_1 \\ w_2^T x - b_2 \\ \vdots \\ w_K^T x - b_K \end{bmatrix}$$

Predict via Largest Score:

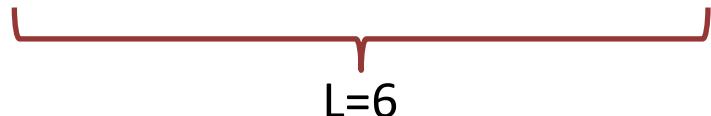
$$\operatorname{argmax}_k \begin{bmatrix} w_1^T x - b_1 \\ w_2^T x - b_2 \\ \vdots \\ w_K^T x - b_K \end{bmatrix}$$

- How many classes?

# Multiclass Prediction

- $x = \text{"Fish Sleep"}$
- $y = (N, V)$
- Multiclass prediction:
  - All possible length-M sequences as different class
  - (D, D), (D, N), (D, V), (D, Adj), (D, Adv), (D, Pr)  
(N, D), (N, N), (N, V), (N, Adj), (N, Adv), ...
- **$L^M$  classes!**
  - Length 2:  $6^2 = 36!$

POS Tags:  
Det, Noun, Verb, Adj, Adv, Prep



$L=6$

# Multiclass Prediction

- $x = \text{"Fish Sleep"}$

POS Tags:  
Det, Noun, Verb, Adj, Adv, Prep

- $y = (N, V)$

- M

Exponential Explosion in #Classes!  
(Not Tractable for Sequence Prediction)

- $L^m$  classes:

- Length 2:  $6^2 = 36!$

# Why is Naïve Multiclass Intractable?

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

- (D, D, D), (D, D, N), (D, D, V), (D, D, Adj), (D, D, Adv), (D, D, Pr)
- (D, N, D), (D, N, N), (D, N, V), (D, N, Adj), (D, N, Adv), (D, N, Pr)
- (D, V, D), (D, V, N), (D, V, V), (D, V, Adj), (D, V, Adv), (D, V, Pr)
- ...
- (N, D, D), (N, D, N), (N, D, V), (N, D, Adj), (N, D, Adv), (N, D, Pr)
- (N, N, D), (N, N, N), (N, N, V), (N, N, Adj), (N, N, Adv), (N, N, Pr)
- ...

Assume pronouns are nouns for simplicity.

# Why is Naïve Multiclass Intractable?

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

- (D, D, D), (D, D, N), (D, D, V), (D, D, Adj), (D, D, Adv), (D, D, Pr)

Treats Every Combination As Different Class  
(Learn model for each combination)

Exponentially Large Representation!  
(Exponential Time to Consider Every Class)  
(Exponential Storage)

Assume pronouns are nouns for simplicity.

# Independent Classification

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

- Treat each word independently (assumption)
  - Independent multiclass prediction per word
  - Predict for  $x = \text{"I"}$  independently
  - Predict for  $x = \text{"fish"}$  independently
  - Predict for  $x = \text{"often"}$  independently
  - Concatenate predictions.

Assume pronouns are nouns for simplicity.

# Independent Classification

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

- Treat each word independently (assumption)
  - Independent multiclass prediction per word

#Classes = #POS Tags  
(6 in our example)

Solvable using standard multiclass prediction.

# Independent Classification

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

- Treat each word independently
  - Independent multiclass prediction per word

$P(y x)$	$x = \text{"I"}$	$x = \text{"fish"}$	$x = \text{"often"}$
y="Det"	0.0	0.0	0.0
y="Noun"	1.0	0.75	0.0
y="Verb"	0.0	0.25	0.0
y="Adj"	0.0	0.0	0.4
y="Adv"	0.0	0.0	0.6
y="Prep"	0.0	0.0	0.0

**Prediction:** (N, N, Adv)

**Correct:** (N, V, Adv)

**Why the mistake?**

Assume pronouns are nouns for simplicity.

# Context Between Words

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

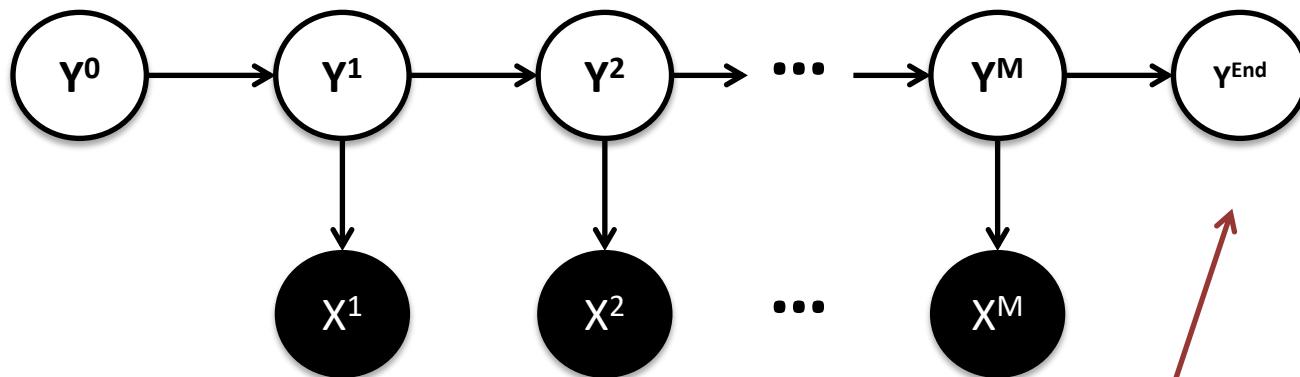
- Independent Predictions Ignore Word Pairs
  - In Isolation:
    - “Fish” is more likely to be a Noun
  - But Conditioned on Following a (pro)Noun...
    - “Fish” is more likely to be a Verb!
  - **“1st Order” Dependence (Model All Pairs)**
    - 2<sup>nd</sup> Order Considers All Triplets
    - Arbitrary Order = Exponential Size (Naïve Multiclass)

Assume pronouns are nouns for simplicity.

# 1<sup>st</sup> Order Hidden Markov Model

- $x = (x^1, x^2, x^3, x^4, \dots, x^M)$  (sequence of words)
- $y = (y^1, y^2, y^3, y^4, \dots, y^M)$  (sequence of POS tags)
- $P(x^i | y^i)$  Probability of state  $y^i$  generating  $x^i$
- $P(y^{i+1} | y^i)$  Probability of state  $y^i$  transitioning to  $y^{i+1}$
- $P(y^1 | y^0)$   $y^0$  is defined to be the Start state
- $P(\text{End} | y^M)$  Prior probability of  $y^M$  being the final state
  - Not always used

# Graphical Model Representation



$$P(x, y) = P(End \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

# 1<sup>st</sup> Order Hidden Markov Model

$$P(x, y) = P(\text{End} | y^M) \prod_{i=1}^M P(y^i | y^{i-1}) \prod_{i=1}^M P(x^i | y^i)$$

“Joint Distribution”

Optional

- $P(x^i | y^i)$  Probability of state  $y^i$  generating  $x^i$
- $P(y^{i+1} | y^i)$  Probability of state  $y^i$  transitioning to  $y^{i+1}$
- $P(y^1 | y^0)$   $y^0$  is defined to be the Start state
- $P(\text{End} | y^M)$  Prior probability of  $y^M$  being the final state
  - Not always used

# 1<sup>st</sup> Order Hidden Markov Model

$$P(x|y) = \prod_{i=1}^M P(x^i | y^i)$$

Given a POS Tag Sequence  $y$ :  
**Can compute each  $P(x^i|y)$  independently!**  
( $x^i$  conditionally independent given  $y^i$ )

“Conditional Distribution on  $x$  given  $y$ ”

- $P(x^i | y^i)$  Probability of state  $y^i$  generating  $x^i$
- $P(y^{i+1} | y^i)$  Probability of state  $y^i$  transitioning to  $y^{i+1}$
- $P(y^1 | y^0)$   $y^0$  is defined to be the Start state
- $P(\text{End} | y^M)$  Prior probability of  $y^M$  being the final state
  - Not always used

# 1<sup>st</sup> Order Hidden Markov Model

## Models All State-State Pairs

(all POS Tag-Tag pairs)

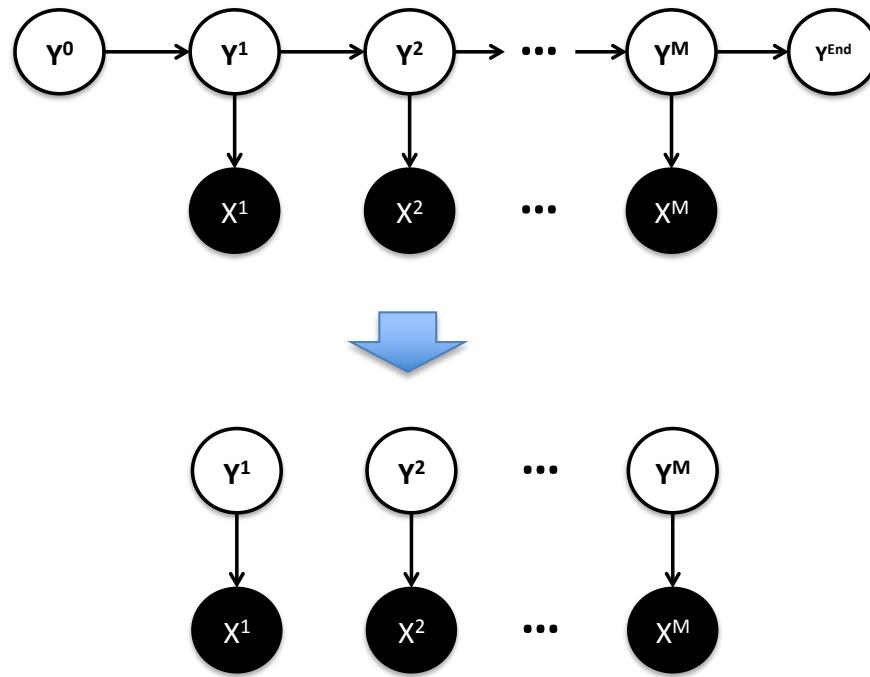
# Models All State-Observation Pairs

(all Tag-Word pairs)

Same Complexity as Independent Multiclass

- $P(x^i | y^i)$  Probability of state  $y^i$  generating  $x^i$
  - $P(y^{i+1} | y^i)$  Probability of state  $y^i$  transitioning to  $y^{i+1}$
  - $P(y^1 | y^0)$   $y^0$  is defined to be the Start state
  - $P(\text{End} | y^M)$  Prior probability of  $y^M$  being the final state
    - Not always used

# Relationship to Naïve Bayes



Reduces to a sequence of disjoint Naïve Bayes models  
(if we ignore transition probabilities)

# $P(\text{word} \mid \text{state/tag})$

- Two-word language: “fish” and “sleep”
- Two-tag language: “Noun” and “Verb”

$P(x y)$	$y = \text{“Noun”}$	$y = \text{“Verb”}$
x=“fish”	0.8	0.5
x=“sleep”	0.2	0.5

**Given Tag Sequence y:**

$$P(\text{“fish sleep”} \mid (\text{N,V})) = 0.8 * 0.5$$

$$P(\text{“fish fish”} \mid (\text{N,V})) = 0.8 * 0.5$$

$$P(\text{“sleep fish”} \mid (\text{V,V})) = 0.8 * 0.5$$

$$P(\text{“sleep sleep”} \mid (\text{N,N})) = 0.2 * 0.2$$

# Sampling

- HMMs are “generative” models
  - Models joint distribution  $P(x,y)$
  - Can generate samples from this distribution
  - First consider conditional distribution  $P(x|y)$

$P(x y)$	$y = \text{“Noun”}$	$y = \text{“Verb”}$
x = “fish”	0.8	0.5
x = “sleep”	0.2	0.5

**Given Tag Sequence  $y = (N,V)$ :**

**Sample each word independently:**

Sample  $P(x^1 | N)$  (0.8 Fish, 0.2 Sleep)

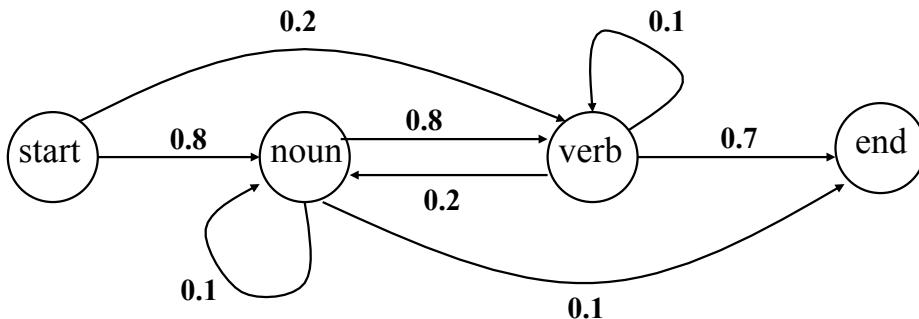
Sample  $P(x^2 | V)$  (0.5 Fish, 0.5 Sleep)

**– What about sampling from  $P(x,y)$ ?**

# Forward Sampling of $P(y|x)$

$$P(x,y) = P(End | y^M) \prod_{i=1}^M P(y^i | y^{i-1}) \prod_{i=1}^M P(x^i | y^i)$$

$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5



Initialize  $y^0 = \text{Start}$

Initialize  $i = 0$

1.  $i = i + 1$
2. Sample  $y^i$  from  $P(y^i | y^{i-1})$
3. If  $y^i == \text{End}$ : Quit
4. Sample  $x^i$  from  $P(x^i | y^i)$
5. Goto Step 1

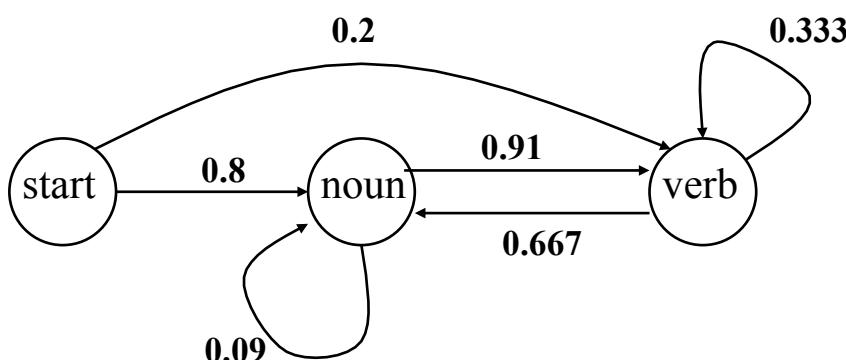
Exploits Conditional Ind.  
Requires  $P(\text{End} | y^i)$

# Forward Sampling of $P(y, x | L)$

$$P(x, y | M) = P(\text{End} | y^M) \prod_{i=1}^M P(y^i | y^{i-1}) \prod_{i=1}^M P(x^i | y^i)$$

~~$P(\text{End} | y^M)$~~

$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5



```

Initialize  $y^0 = \text{Start}$ 
Initialize  $i = 0$ 

1.  $i = i + 1$ 
2. If( $i == M$ ): Quit
3. Sample  $y^i$  from  $P(y^i | y^{i-1})$ 
4. Sample  $x^i$  from  $P(x^i | y^i)$ 
5. Goto Step 1
  
```

Exploits Conditional Ind.  
Assumes no  $P(\text{End} | y^i)$

# 1<sup>st</sup> Order Hidden Markov Model

$$P(x^{k+1:M}, y^{k+1:M} | x^{1:k}, y^{1:k}) = P(x^{k+1:M}, y^{k+1:M} | y^k)$$

“Memory-less Model” – only needs  $y^k$  to model rest of sequence

- $P(x^i | y^i)$  Probability of state  $y^i$  generating  $x^i$
- $P(y^{i+1} | y^i)$  Probability of state  $y^i$  transitioning to  $y^{i+1}$
- $P(y^1 | y^0)$   $y^0$  is defined to be the Start state
- $P(\text{End} | y^M)$  Prior probability of  $y^M$  being the final state
  - Not always used

# Viterbi Algorithm

# Most Common Prediction Problem

- Given input sentence, predict POS Tag seq.

$$\underset{y}{\operatorname{argmax}} P(y | x)$$

- Naïve approach:**
  - Try all possible  $y$ 's
  - Choose one with highest probability
  - Exponential time:  $L^M$  possible  $y$ 's**

# Bayes's Rule

$$\begin{aligned}\operatorname{argmax}_y P(y \mid x) &= \operatorname{argmax}_y \frac{P(y, x)}{P(x)} \\&= \operatorname{argmax}_y P(y, x) \\&= \operatorname{argmax}_y P(x \mid y)P(y)\end{aligned}$$

$$P(x|y) = \prod_{i=1}^M P(x^i|y^i)$$

$$P(y) = P(END|y^M) \prod_{i=1}^M P(y^i|y^{i-1})$$

$$\begin{aligned}
\underset{y}{\operatorname{argmax}} P(y, x) &= \underset{y}{\operatorname{argmax}} \prod_{i=1}^M P(y^i | y^{i-1}) \prod_{i=1}^M P(x^i | y^i) \\
&= \underset{y^M}{\operatorname{argmax}} \underset{y^{1:M-1}}{\operatorname{argmax}} \prod_{i=1}^M P(y^i | y^{i-1}) \prod_{i=1}^M P(x^i | y^i) \\
&= \underset{y^M}{\operatorname{argmax}} \underset{y^{1:M-1}}{\operatorname{argmax}} P(y^M | y^{M-1}) P(x^M | y^M) P(y^{1:M-1} | x^{1:M-1})
\end{aligned}$$

$$\begin{aligned}
P(y^{1:k} | x^{1:k}) &= P(x^{1:k} | y^{1:k}) P(y^{1:k}) & P(x^{1:k} | y^{1:k}) &= \prod_{i=1}^k P(x^i | y^i) \\
P(y^{1:k}) &= \prod_{i=1}^k P(y^{i+1} | y^i)
\end{aligned}$$

**Exploit Memory-less Property:**  
**The choice of  $y^M$  only depends on  $y^{1:M-1}$  via  $P(y^M | y^{M-1})!$**

# Dynamic Programming

- **Input:**  $x = (x^1, x^2, x^3, \dots, x^M)$
- **Computed:** best length- $k$  prefix ending in each Tag:
  - Examples:

$$\hat{Y}^k(V) = \left( \underset{y^{1:k-1}}{\operatorname{argmax}} P(y^{1:k-1} \oplus V, x^{1:k}) \right) \oplus V \quad \hat{Y}^k(N) = \left( \underset{y^{1:k-1}}{\operatorname{argmax}} P(y^{1:k-1} \oplus N, x^{1:k}) \right) \oplus N$$

Sequence Concatenation

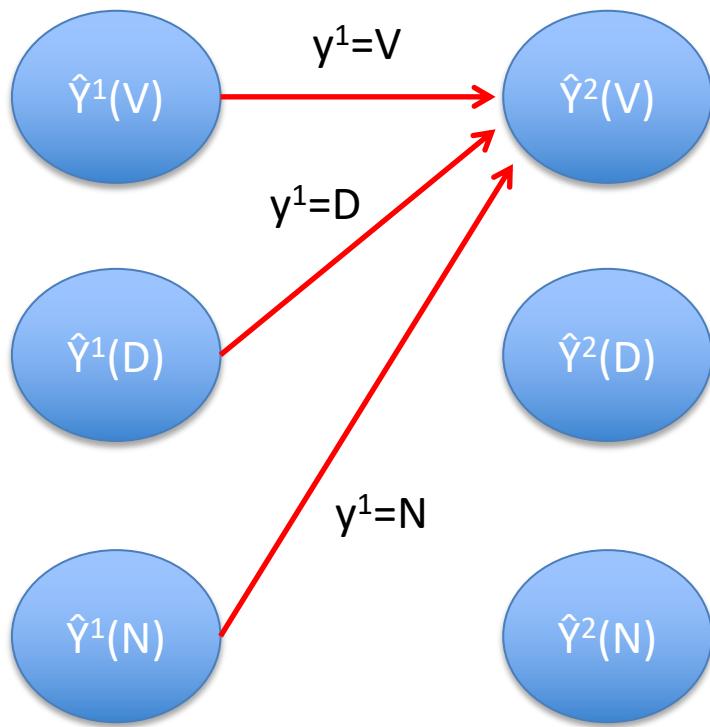
- **Claim:**  $\hat{Y}^{k+1}(V) = \left( \underset{y^{1:k} \in \{\hat{Y}^k(T)\}_T}{\operatorname{argmax}} P(y^{1:k} \oplus V, x^{1:k+1}) \right) \oplus V$

Pre-computed

Recursive Definition!

**Solve:**  $\hat{Y}^2(V) = \left( \underset{y^1 \in \{\hat{Y}^1(T)\}_T}{\operatorname{argmax}} P(y^1, x^1) P(y^2 = V | y^1) P(x^2 | y^2 = V) \right) \oplus V$

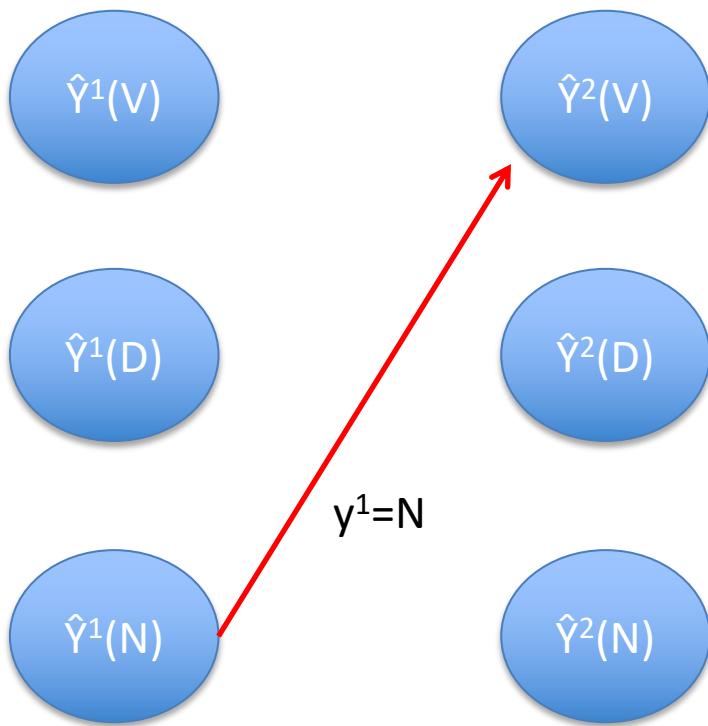
Store each  
 $\hat{Y}^1(Z)$  &  $P(\hat{Y}^1(Z), x^1)$



$\hat{Y}^1(Z)$  is just Z

**Solve:**  $\hat{Y}^2(V) = \left( \underset{y^1 \in \{\hat{Y}^1(T)\}_T}{\operatorname{argmax}} P(y^1, x^1) P(y^2 = V | y^1) P(x^2 | y^2 = V) \right) \oplus V$

Store each  
 $\hat{Y}^1(Z)$  &  $P(\hat{Y}^1(Z), x^1)$



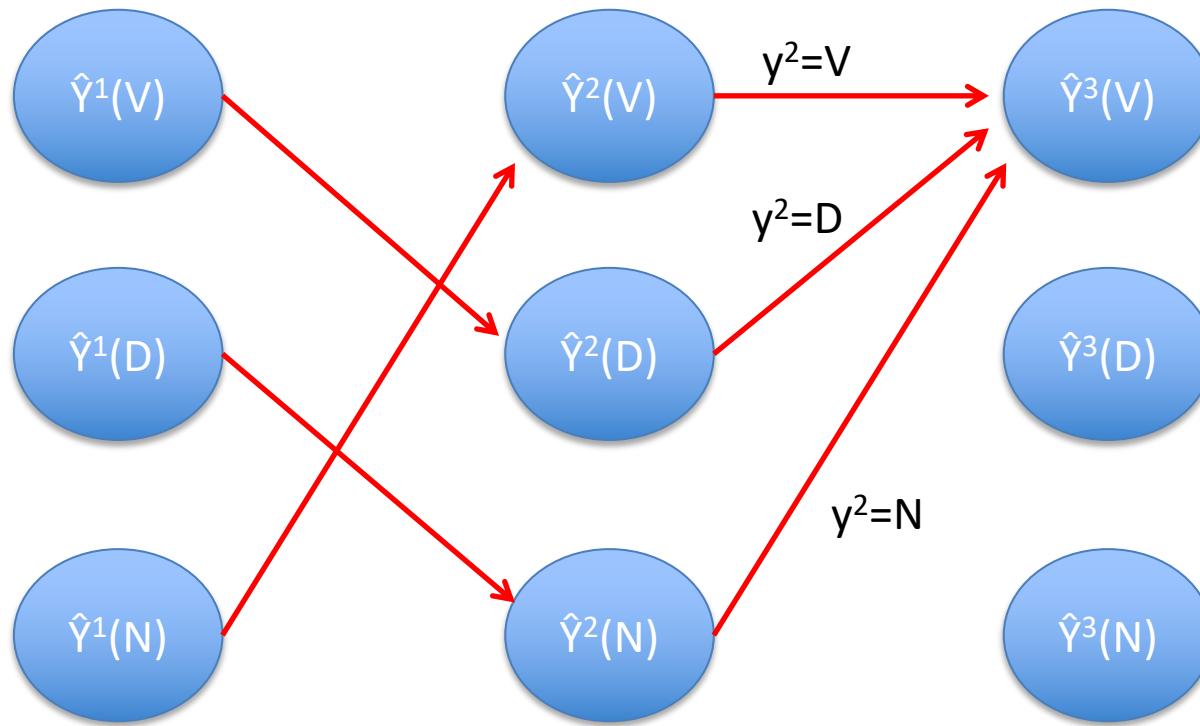
$\hat{Y}^1(Z)$  is just  $Z$

Ex:  $\hat{Y}^2(V) = (N, V)$

**Solve:**  $\hat{Y}^3(V) = \left( \underset{y^{1:2} \in \{\hat{Y}^2(T)\}_T}{\operatorname{argmax}} P(y^{1:2}, x^{1:2}) P(y^3 = V | y^2) P(x^3 | y^3 = V) \right) \oplus V$

Store each  
 $\hat{Y}^1(Z)$  &  $P(\hat{Y}^1(Z), x^1)$

Store each  
 $\hat{Y}^2(Z)$  &  $P(\hat{Y}^2(Z), x^{1:2})$



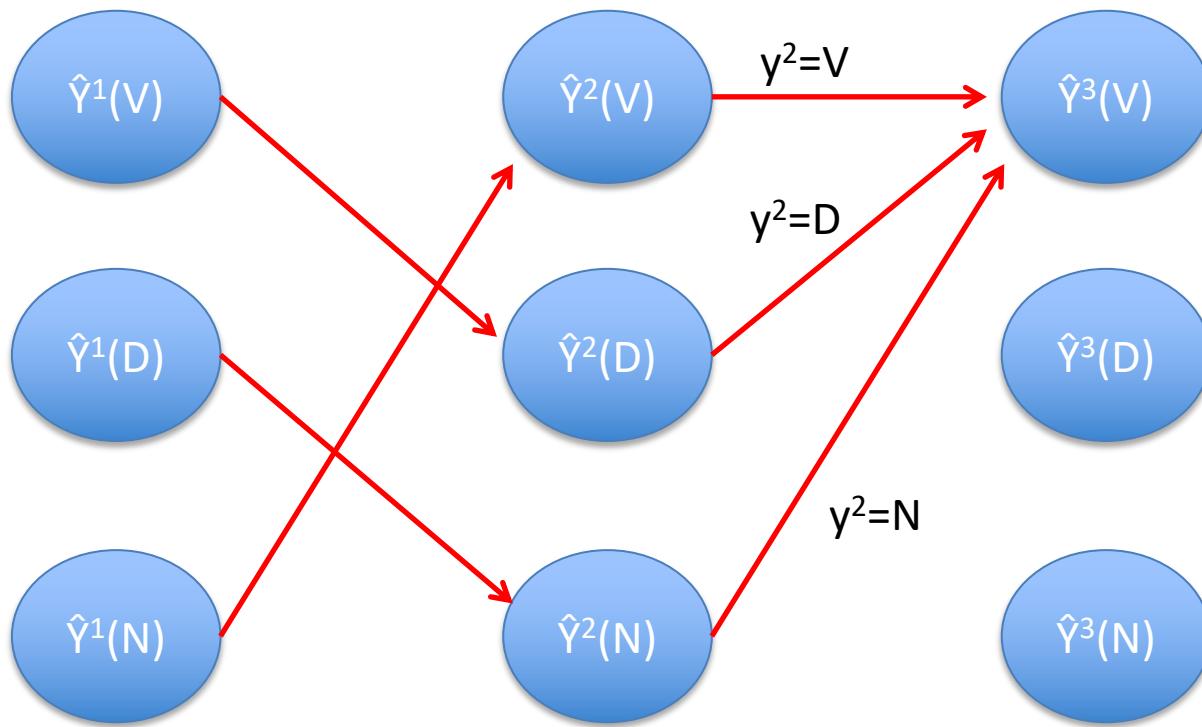
Ex:  $\hat{Y}^2(V) = (N, V)$

**Solve:**  $\hat{Y}^3(V) = \left( \underset{y^{1:2} \in \{\hat{Y}^2(T)\}_T}{\operatorname{argmax}} P(y^{1:2}, x^{1:2}) P(y^3 = V | y^2) P(x^3 | y^3 = V) \right) \oplus V$

Store each  
 $\hat{Y}^1(Z)$  &  $P(\hat{Y}^1(Z), x^1)$

Store each  
 $\hat{Y}^2(Z)$  &  $P(\hat{Y}^2(Z), x^{1:2})$

**Claim:** Only need to check  
solutions of  $\hat{Y}^2(Z)$ ,  $Z=V,D,N$



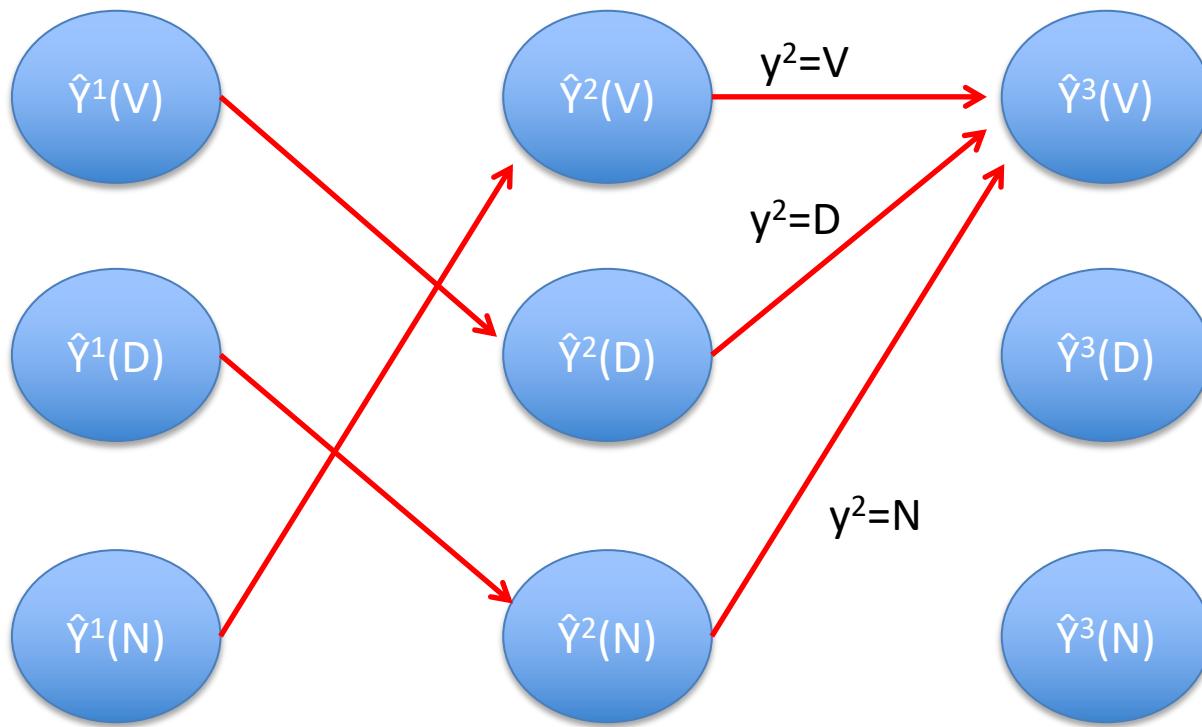
Ex:  $\hat{Y}^2(V) = (N, V)$

**Solve:**  $\hat{Y}^3(V) = \left( \underset{y^{1:2} \in \{\hat{Y}^2(T)\}_T}{\operatorname{argmax}} P(y^{1:2}, x^{1:2}) P(y^3 = V | y^2) P(x^3 | y^3 = V) \right) \oplus V$

Store each  
 $\hat{Y}^1(Z)$  &  $P(\hat{Y}^1(Z), x^1)$

Store each  
 $\hat{Y}^2(Z)$  &  $P(\hat{Y}^2(Z), x^{1:2})$

**Claim:** Only need to check  
solutions of  $\hat{Y}^2(Z)$ ,  $Z=V,D,N$



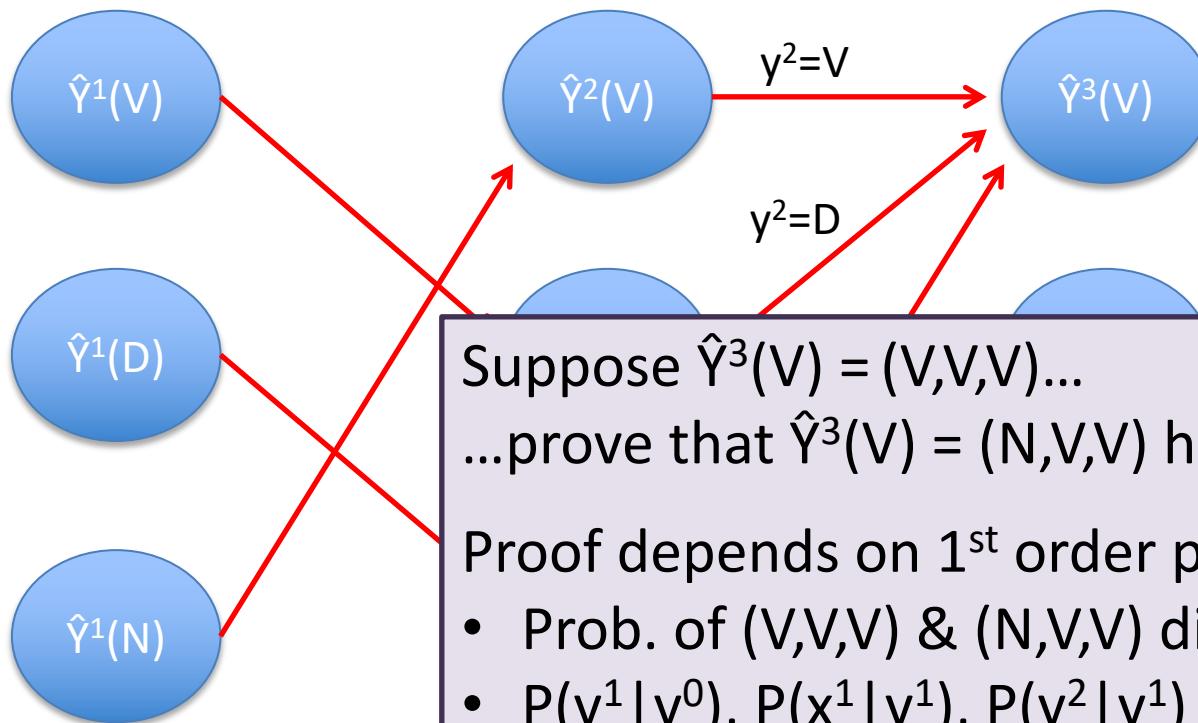
Ex:  $\hat{Y}^2(V) = (N, V)$

**Solve:**  $\hat{Y}^3(V) = \left( \underset{y^{1:2} \in \{\hat{Y}^2(T)\}_T}{\operatorname{argmax}} P(y^{1:2}, x^{1:2}) P(y^3 = V | y^2) P(x^3 | y^3 = V) \right) \oplus V$

Store each  
 $\hat{Y}^1(Z)$  &  $P(\hat{Y}^1(Z), x^1)$

Store each  
 $\hat{Y}^2(Z)$  &  $P(\hat{Y}^2(Z), x^{1:2})$

**Claim:** Only need to check  
solutions of  $\hat{Y}^2(Z)$ ,  $Z=V,D,N$



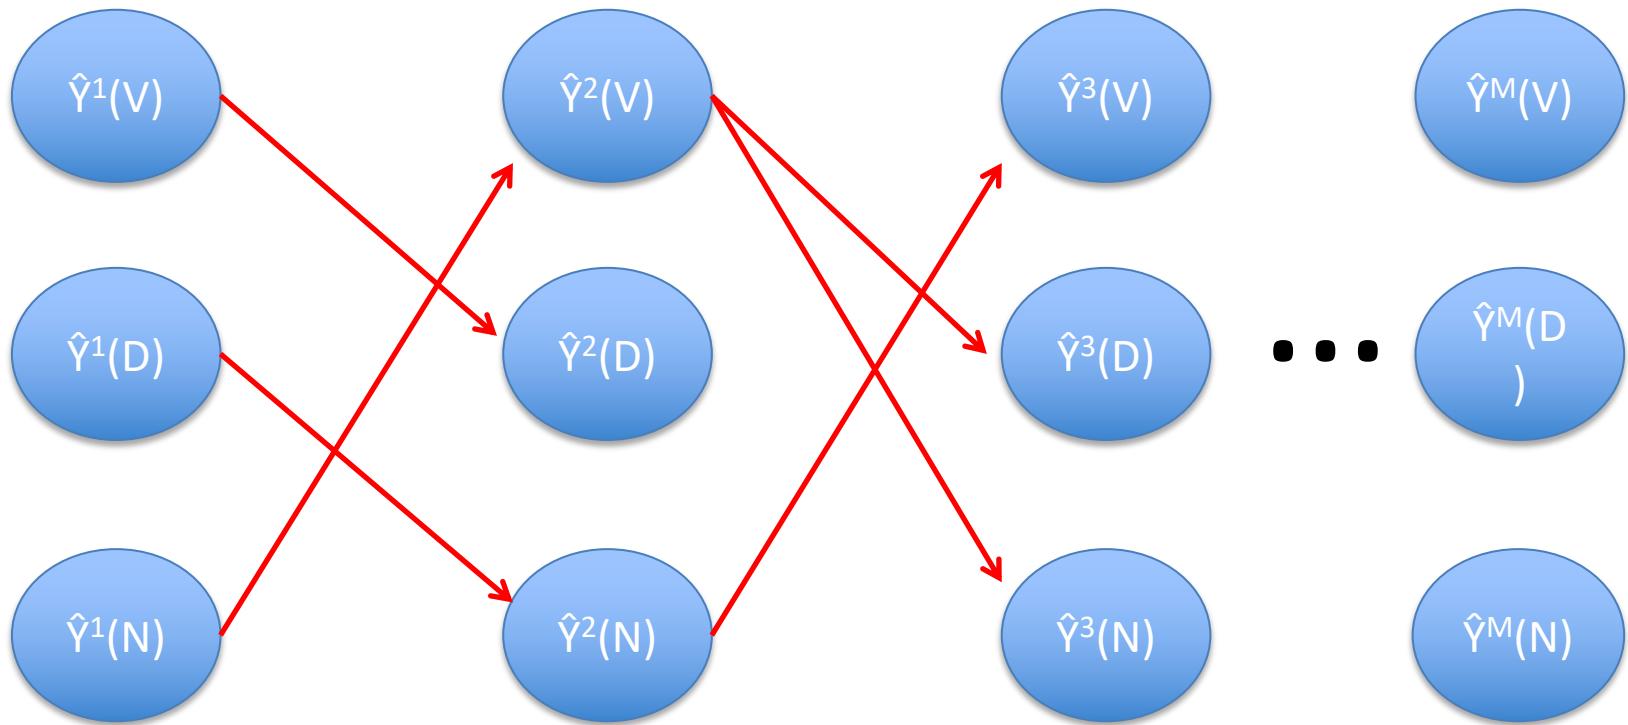
$$\hat{Y}^M(V) = \left( \underset{y^{1:M-1} \in \{\hat{Y}^{M-1}(T)\}_T}{\operatorname{argmax}} P(y^{1:M-1}, x^{1:M-1}) P(y^M = V \mid y^{M-1}) P(x^M \mid y^M = V) P(\text{End} \mid y^M = V) \right) \oplus V$$

Optional

Store each  
 $\hat{Y}^1(Z) \text{ & } P(\hat{Y}^1(Z), x^1)$

Store each  
 $\hat{Y}^2(Z) \text{ & } P(\hat{Y}^2(Z), x^{1:2})$

Store each  
 $\hat{Y}^3(Z) \text{ & } P(\hat{Y}^3(Z), x^{1:3})$



Ex:  $\hat{Y}^2(V) = (N, V)$

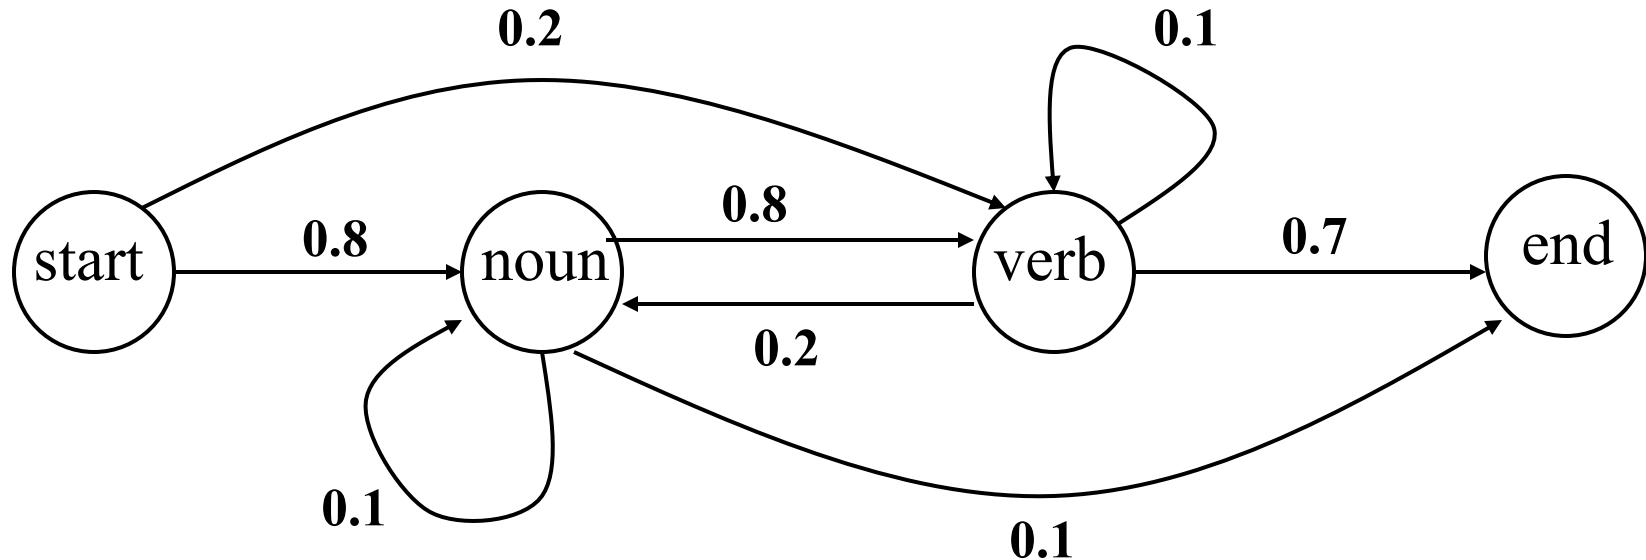
Ex:  $\hat{Y}^3(V) = (D, N, V)$

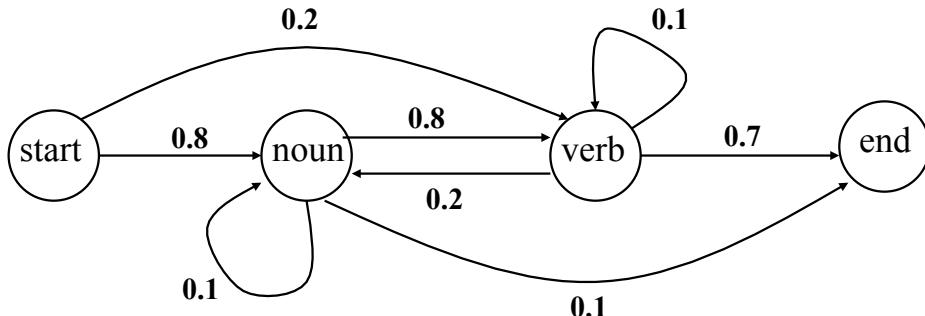
# Viterbi Algorithm

- Solve:
$$\begin{aligned}\operatorname{argmax}_y P(y \mid x) &= \operatorname{argmax}_y \frac{P(y, x)}{P(x)} \\ &= \operatorname{argmax}_y P(y, x) \\ &= \operatorname{argmax}_y P(x \mid y)P(y)\end{aligned}$$
- For  $k=1..M$ 
  - Iteratively solve for each  $\hat{Y}^k(Z)$ 
    - $Z$  looping over every POS tag.
- Predict best  $\hat{Y}^M(Z)$
- Also known as Mean A Posteriori (MAP) inference

# Numerical Example

$x = (\text{Fish Sleep})$





$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

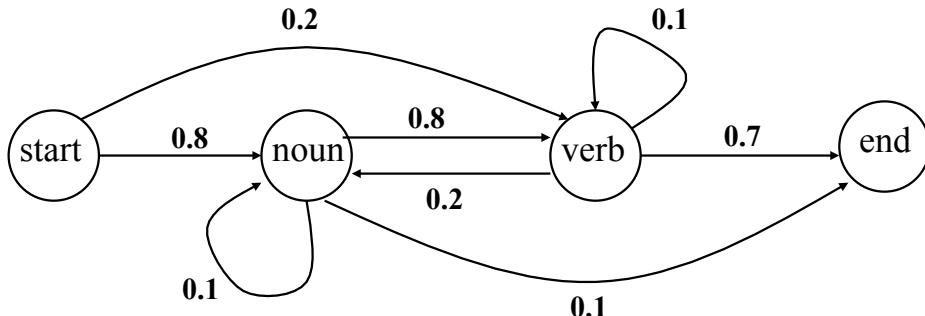
0                  1                  2                  3

start              1

verb              0

noun              0

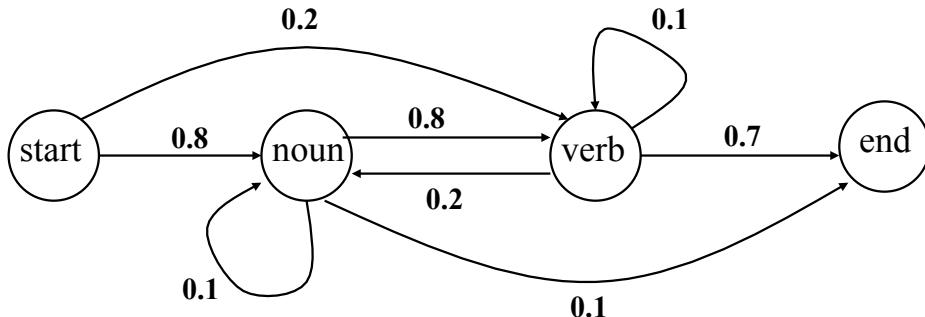
end                0



$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

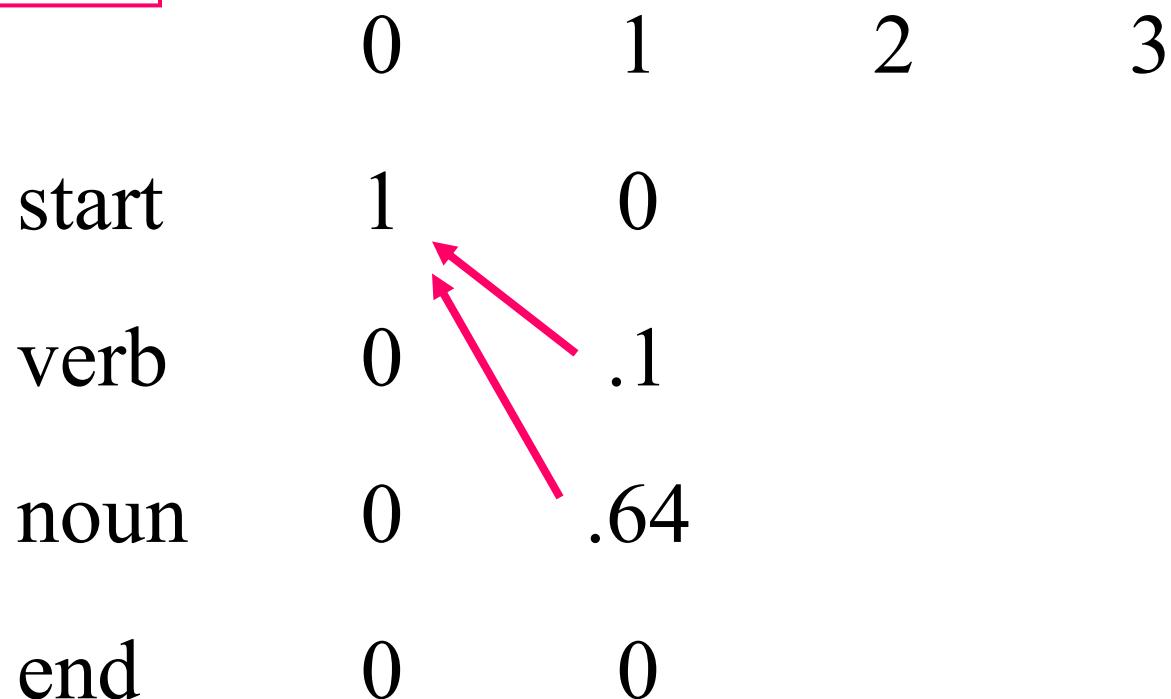
Token 1: fish

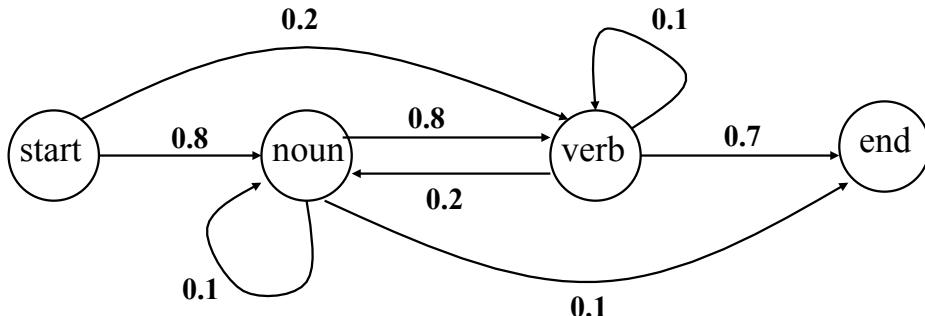
	0	1	2	3
start	1	0		
verb	0	.2 * .5		
noun	0	.8 * .8		
end	0	0		



$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 1: fish





$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 2: sleep

(if 'fish' is verb)

0            1            2            3

start

1            0            0

verb

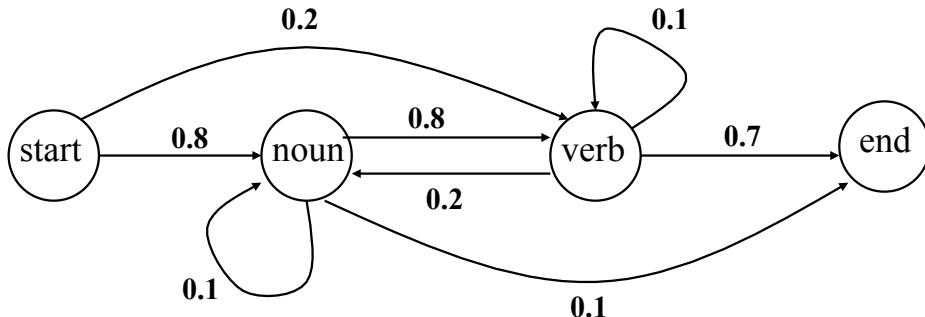
0            .1             $.1 * .1 * .5$

noun

0             $.64$              $.1 * .2 * .2$

end

0            0            -



$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 2: sleep

(if 'fish' is verb)

0            1            2            3

start

1            0            0

verb

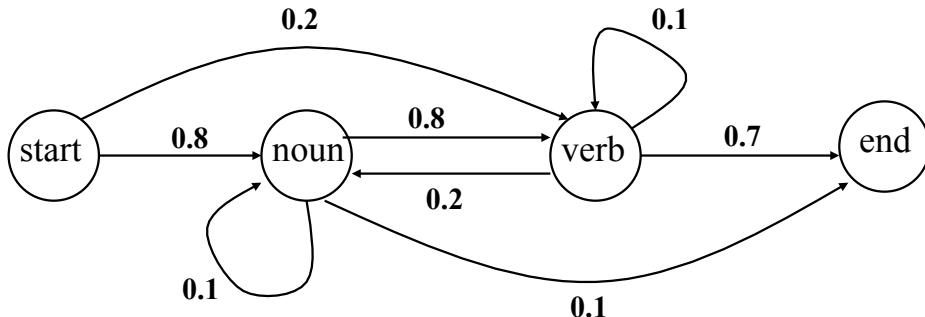
0            .1            .005

noun

0            .64            .004

end

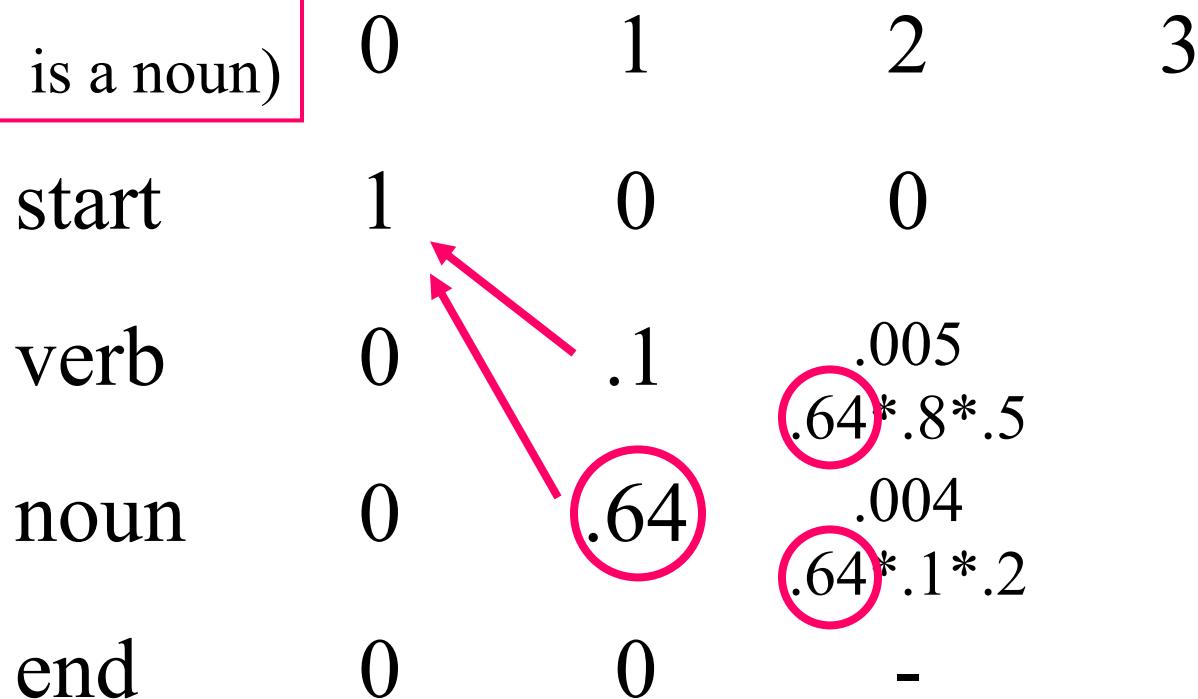
0            0            -

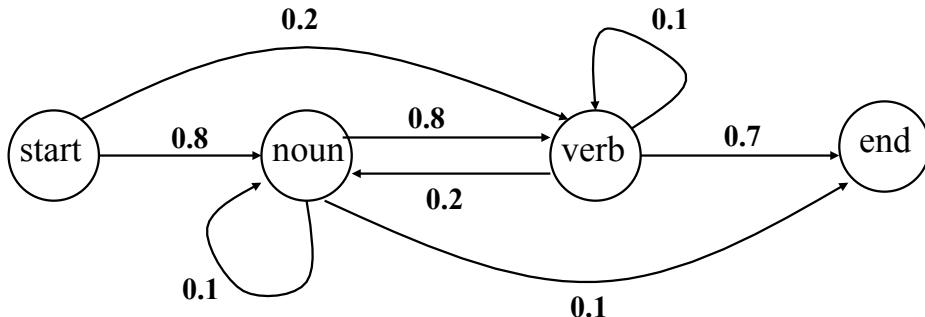


$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 2: sleep

(if 'fish' is a noun)



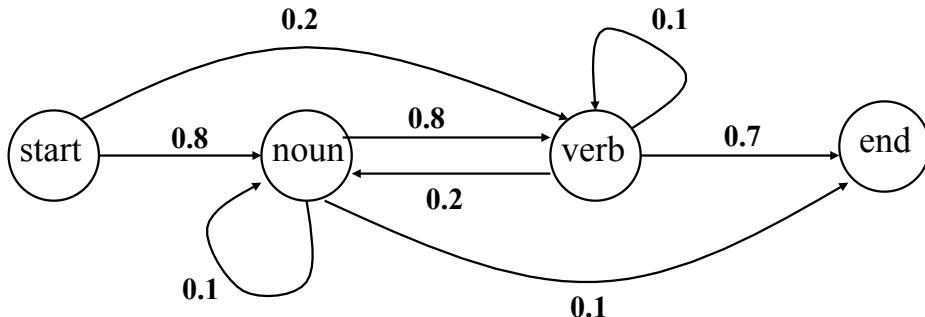


$P(x y)$	$y = \text{Noun}$	$y = \text{Verb}$
$x = \text{fish}$	0.8	0.5
$x = \text{sleep}$	0.2	0.5

Token 2: sleep

(if 'fish' is a noun)

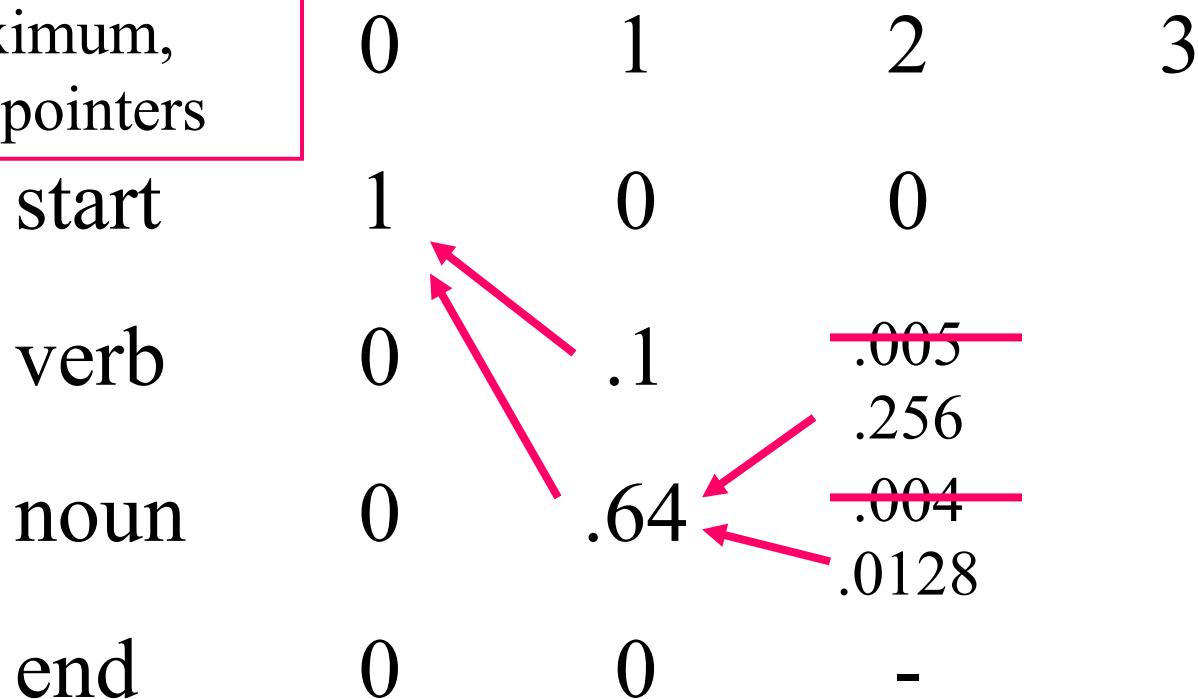
	0	1	2	3
start	1	0	0	
verb	0	.1	.005	.256
noun	0	.64	.004	.0128
end	0	0	-	

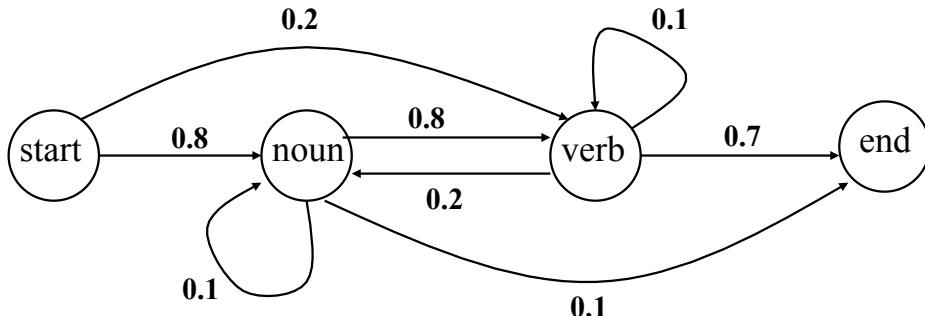


$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 2: sleep

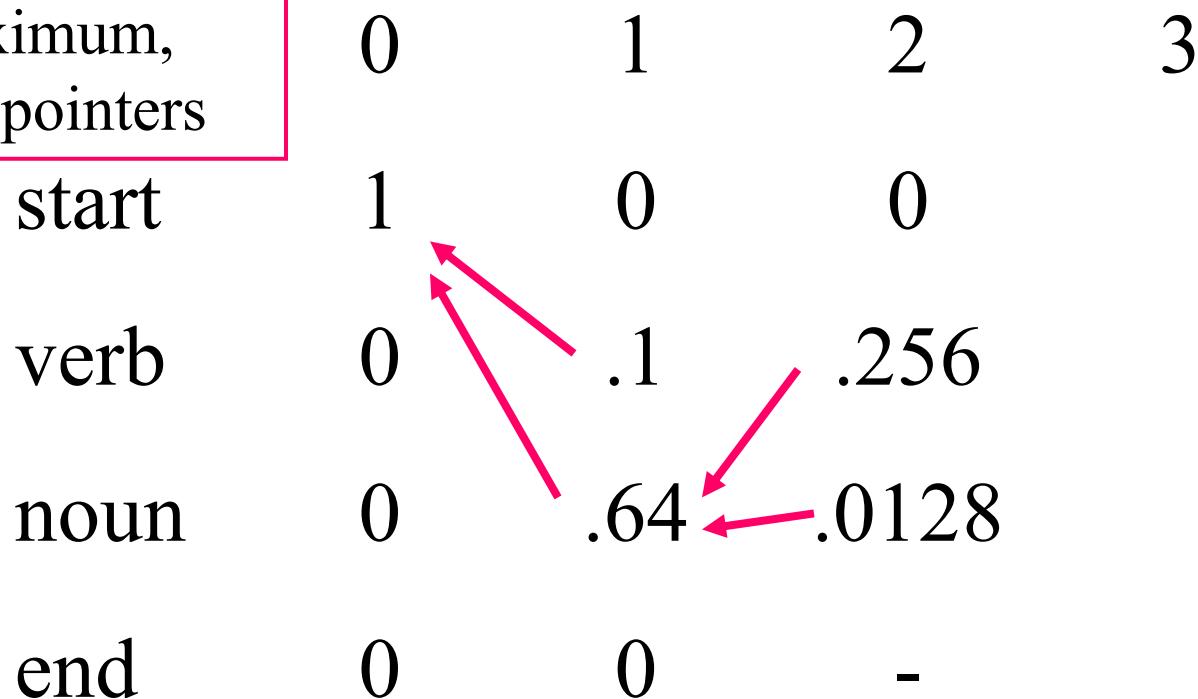
take maximum,  
set back pointers

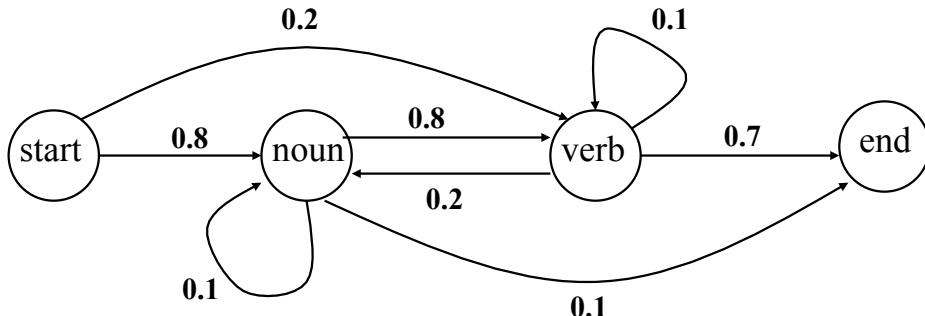




$P(x y)$	$y = \text{Noun}$	$y = \text{Verb}$
$x = \text{fish}$	0.8	0.5
$x = \text{sleep}$	0.2	0.5

Token 2: sleep  
 take maximum,  
 set back pointers

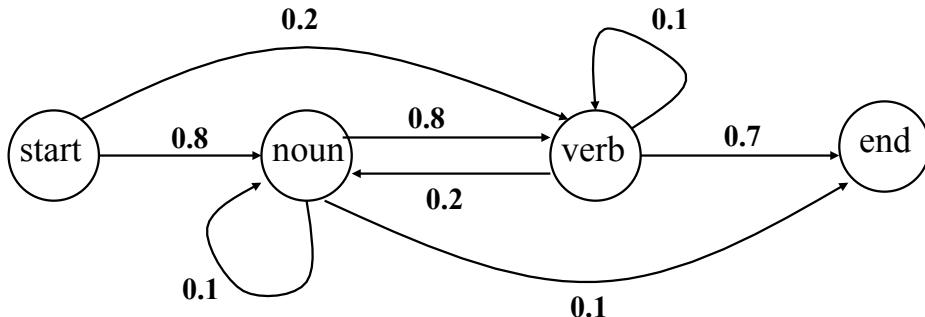




$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 3: end

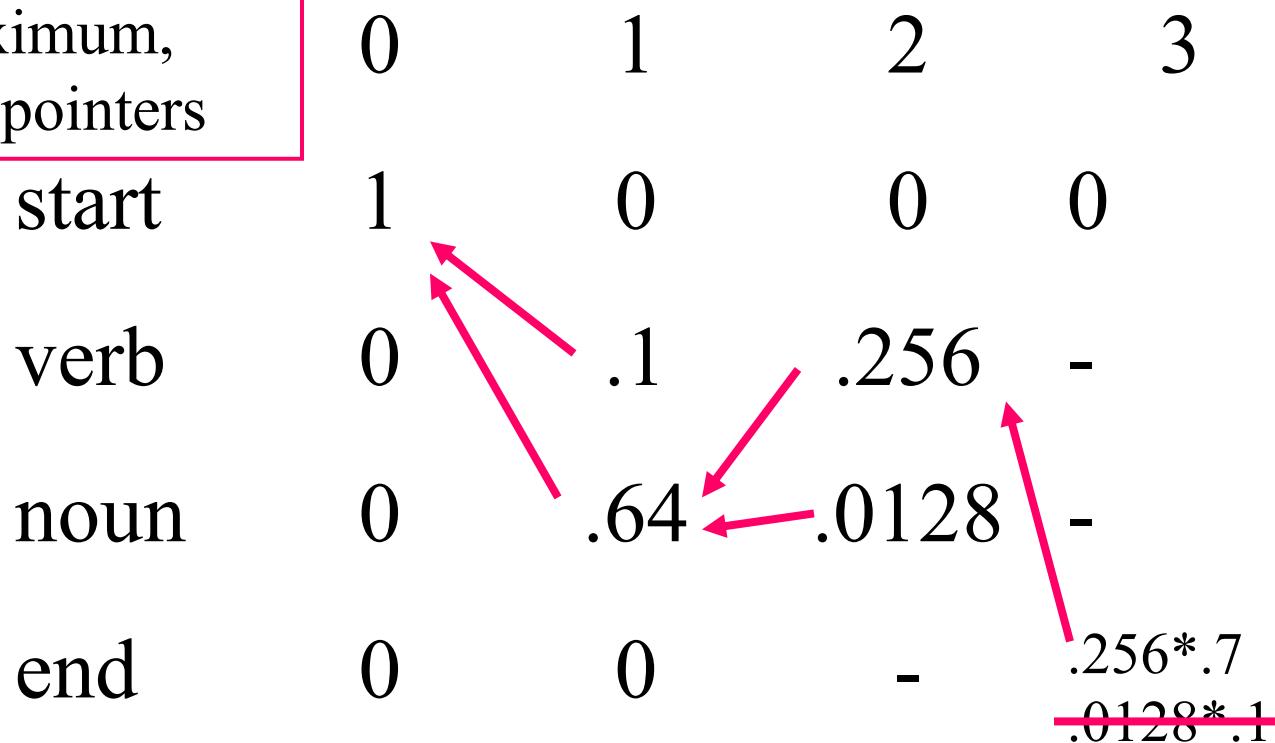
	0	1	2	3
start	1	0	0	0
verb	0	.1	.256	-
noun	0	.64	.0128	-
end	0	0	-	$.256 * .7$ $.0128 * .1$

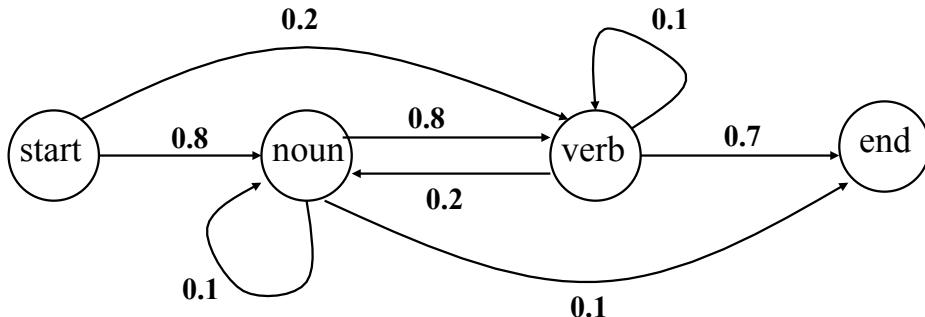


$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

Token 3: end

take maximum,  
set back pointers



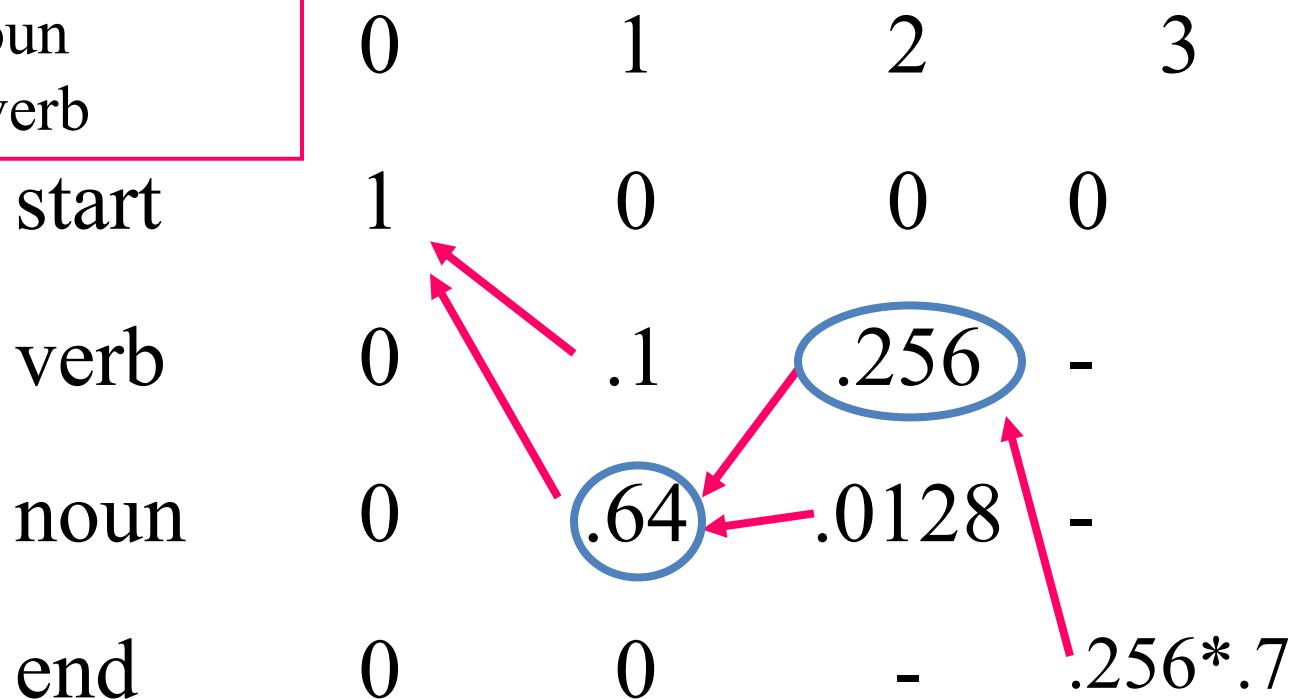


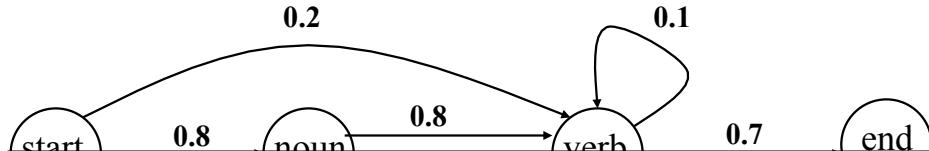
$P(x y)$	$y = \text{Noun}$	$y = \text{Verb}$
$x = \text{fish}$	0.8	0.5
$x = \text{sleep}$	0.2	0.5

Decode:

fish = noun

sleep = verb





$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5

## What might go wrong for long sequences?

fish = noun

sleep = verb

start



## Underflow!

Small numbers get repeatedly multiplied together – exponentially small!

end



# Viterbi Algorithm (w/ Log Probabilities)

- Solve:  $\underset{y}{\operatorname{argmax}} P(y \mid x) = \underset{y}{\operatorname{argmax}} \frac{P(y, x)}{P(x)}$   
 $= \underset{y}{\operatorname{argmax}} P(y, x)$   
 $= \underset{y}{\operatorname{argmax}} \log P(x \mid y) + \log P(y)$
- For  $k=1..M$ 
  - Iteratively solve for each  $\hat{Y}^k(Z)$ 
    - Z looping over every POS tag.
- Predict best  $\hat{Y}^M(Z)$ 
  - **Log( $\hat{Y}^M(Z)$ ) accumulates additively, not multiplicatively**

# Recap: Independent Classification

$x = \text{"I fish often"}$

**POS Tags:**  
Det, Noun, Verb, Adj, Adv, Prep

- Treat each word independently
  - Independent multiclass prediction per word

$P(y x)$	$x = \text{"I"}$	$x = \text{"fish"}$	$x = \text{"often"}$
y="Det"	0.0	0.0	0.0
y="Noun"	1.0	0.75	0.0
y="Verb"	0.0	0.25	0.0
y="Adj"	0.0	0.0	0.4
y="Adv"	0.0	0.0	0.6
y="Prep"	0.0	0.0	0.0

**Prediction:** (N, N, Adv)

**Correct:** (N, V, Adv)

**Mistake due to not  
modeling multiple words.**

Assume pronouns are nouns for simplicity.

# Recap: Viterbi

- Models pairwise transitions between states
  - Pairwise transitions between POS Tags
  - “1<sup>st</sup> order” model

$$P(x, y) = P(End \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

$x = \text{"I fish often"}$       **Independent:** (N, N, Adv)

**HMM Viterbi:** (N, V, Adv)

\*Assuming we defined  $P(x,y)$  properly

# Training HMMs

# Supervised Training

- **Given:**

$$S = \{(x_i, y_i)\}_{i=1}^N$$

Word Sequence (Sentence)      POS Tag Sequence

- **Goal:** Estimate  $P(x, y)$  using  $S$

$$P(x, y) = P(End \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

- **Maximum Likelihood!**

# Aside: Matrix Formulation

- Define Transition Matrix: A

- $A_{ab} = P(y^{i+1}=a | y^i=b)$  or  $-\text{Log}(P(y^{i+1}=a | y^i=b))$

$P(y^{\text{next}}   y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$y^{\text{next}} = \text{"Noun"}$	0.09	0.667
$y^{\text{next}} = \text{"Verb"}$	0.91	0.333

- Observation Matrix: O

- $O_{wz} = P(x^i=w | y^i=z)$  or  $-\text{Log}(P(x^i=w | y^i=z))$

$P(x y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

# Aside: Matrix Formulation

$$P(x, y) = P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

$$P(x, y) = P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

$$= A_{\text{End}, y^M} \prod_{i=1}^M A_{y^i, y^{i-1}} \prod_{i=1}^M O_{x^i, y^i}$$

$$-\log(P(x, y)) = \tilde{A}_{\text{End}, y^M} + \sum_{i=1}^M \tilde{A}_{y^i, y^{i-1}} + \sum_{i=1}^M \tilde{O}_{x^i, y^i}$$

Log prob. formulation  
Each entry of  $\tilde{A}$  is  
defined as  $-\log(A)$

# Maximum Likelihood

$$\operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(x,y) = \operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

- Estimate each component separately:

$$A_{ab} = \frac{\sum_{j=1}^N \sum_{i=0}^{M_j} 1_{[(y_j^{i+1}=a) \wedge (y_j^i=b)]}}{\sum_{j=1}^N \sum_{i=0}^{M_j} 1_{[y_j^i=b]}}$$

$$O_{wz} = \frac{\sum_{j=1}^N \sum_{i=1}^{M_j} 1_{[(x_j^i=w) \wedge (y_j^i=z)]}}{\sum_{j=1}^N \sum_{i=1}^{M_j} 1_{[y_j^i=z]}}$$

- (Derived via minimizing neg. log likelihood)

# Recap: Supervised Training

$$\operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(x,y) = \operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

- Maximum Likelihood Training
  - Counting statistics
  - Super easy!
  - Why?
- What about unsupervised case?

# Recap: Supervised Training

$$\operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(x,y) = \operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

- Maximum Likelihood Training
  - Counting statistics
  - Super easy!
  - Why?
- What about unsupervised case?

# Conditional Independence Assumptions

$$\operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(x,y) = \operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

- Everything decomposes to products of pairs
  - I.e.,  $P(y^{i+1}=a \mid y^i=b)$  doesn't depend on anything else
- Can just estimate frequencies:
  - How often  $y^{i+1}=a$  when  $y^i=b$  over training set
  - Note that  $P(y^{i+1}=a \mid y^i=b)$  is a common model across all locations of all sequences.

# Conditional Independence Assumptions

$$\operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(x,y) = \operatorname{argmax}_{A,O} \prod_{(x,y) \in S} P(\text{End} \mid y^M) \prod_{i=1}^M P(y^i \mid y^{i-1}) \prod_{i=1}^M P(x^i \mid y^i)$$

# Parameters:  
Transitions A: #Tags<sup>2</sup>  
Observations O: #Words x #Tags

Avoids directly model word/word pairings  
#Tags = 10s  
#Words = 10000s

# Unsupervised Training

- What about if no y's?
  - Just a training set of sentences

$$S = \{x_i\}_{i=1}^N$$

  
Word Sequence  
(Sentence)

- Still want to estimate  $P(x,y)$ 
  - How?
  - Why?

$$\arg \max \prod_i P(x_i) = \operatorname{argmax} \prod_i \sum_y P(x_i, y)$$

# Unsupervised Training

- What about if no y's?
  - Just a training set of sentences

$$S = \{x_i\}_{i=1}^N$$

  
Word Sequence  
(Sentence)

- Still want to estimate  $P(x,y)$ 
  - How?
  - Why?

$$\arg \max \prod_i P(x_i) = \operatorname{argmax} \prod_i \sum_y P(x_i, y)$$

# Why Unsupervised Training?

- Supervised Data hard to acquire
  - Require annotating POS tags
- Unsupervised Data plentiful
  - Just grab some text!
- Might just work for POS Tagging!
  - Learn y's that correspond to POS Tags
- Can be used for other tasks
  - Detect outlier sentences (sentences with low prob.)
  - Sampling new sentences.

# EM Algorithm (Baum-Welch)

- If we had  $y$ 's  $\rightarrow$  max likelihood.
- If we had  $(A, O)$   $\rightarrow$  predict  $y$ 's

Chicken vs Egg!

1. Initialize  $A$  and  $O$  arbitrarily
2. Predict prob. of  $y$ 's for each training  $x$
3. Use  $y$ 's to estimate new  $(A, O)$
4. Repeat back to Step 1 until convergence

Expectation Step



Maximization Step



# Expectation Step

- Given  $(A, O)$
- For training  $x = (x^1, \dots, x^M)$ 
  - Predict  $P(y^i)$  for each  $y = (y^1, \dots, y^M)$

	$x^1$	$x^2$	...	$x^L$
$P(y^i = \text{Noun})$	0.5	0.4	...	0.05
$P(y^i = \text{Det})$	0.4	0.6	...	0.25
$P(y^i = \text{Verb})$	0.1	0.0	...	0.7

- Encodes current model's beliefs about  $y$
- “Marginal Distribution” of each  $y^i$

# Recall: Matrix Formulation

- Define Transition Matrix: A
  - $A_{ab} = P(y^{i+1}=a | y^i=b)$  or  $-\text{Log}(P(y^{i+1}=a | y^i=b))$

$P(y^{\text{next}}   y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$y^{\text{next}} = \text{"Noun"}$	0.09	0.667
$y^{\text{next}} = \text{"Verb"}$	0.91	0.333

- Observation Matrix: O
  - $O_{wz} = P(x^i=w | y^i=z)$  or  $-\text{Log}(P(x^i=w | y^i=z))$

$P(x   y)$	$y = \text{"Noun"}$	$y = \text{"Verb"}$
$x = \text{"fish"}$	0.8	0.5
$x = \text{"sleep"}$	0.2	0.5

# Maximization Step

- Max. Likelihood over Marginal Distribution

**Supervised:**  $A_{ab} = \frac{\sum_{j=1}^N \sum_{i=0}^{M_j} 1_{[(y_j^{i+1}=a) \wedge (y_j^i=b)]}}{\sum_{j=1}^N \sum_{i=0}^{M_j} 1_{[y_j^i=b]}}$

Marginals

**Unsupervised:**

$$A_{ab} = \frac{\sum_{j=1}^N \sum_{i=0}^{M_j} P(y_j^i = b, y_j^{i+1} = a)}{\sum_{j=1}^N \sum_{i=0}^{M_j} P(y_j^i = b)}$$

Marginals

$$O_{wz} = \frac{\sum_{j=1}^N \sum_{i=1}^{M_j} 1_{[(x_j^i=w) \wedge (y_j^i=z)]}}{\sum_{j=1}^N \sum_{i=1}^{M_j} 1_{[y_j^i=z]}}$$

Marginals

$$O_{wz} = \frac{\sum_{j=1}^N \sum_{i=1}^{M_j} 1_{[x_j^i=w]} P(y_j^i = z)}{\sum_{j=1}^N \sum_{i=1}^{M_j} P(y_j^i = z)}$$

# Computing Marginals

## (Forward-Backward Algorithm)

- Solving E-Step, requires compute marginals

	$x^1$	$x^2$	...	$x^L$
$P(y^i=\text{Noun})$	0.5	0.4	...	0.05
$P(y^i=\text{Det})$	0.4	0.6	...	0.25
$P(y^i=\text{Verb})$	0.1	0.0	...	0.7

- Can solve using Dynamic Programming!
  - Similar to Viterbi

# Notation

Probability of observing prefix  $x^{1:i}$  and having the i-th state be  $y^i=z$

$$\alpha_z(i) = P(x^{1:i}, y^i = z | A, O)$$

Probability of observing suffix  $x^{i+1:M}$  given the i-th state being  $y^i=z$

$$\beta_z(i) = P(x^{i+1:M} | y^i = z, A, O)$$

Computing Marginals = Combining the Two Terms

$$P(y^i = z | x) = \frac{a_z(i)\beta_z(i)}{\sum_{z'} a_{z'}(i)\beta_{z'}(i)}$$

# Notation

Probability of observing prefix  $x^{1:i}$  and having the i-th state be  $y^i = Z$

$$\alpha_z(i) = P(x^{1:i}, y^i = Z | A, O)$$

Probability of observing suffix  $x^{i+1:M}$  given the i-th state being  $y^i = Z$

$$\beta_z(i) = P(x^{i+1:M} | y^i = Z, A, O)$$

Computing Marginals = Combining the Two Terms

$$P(y^i = b, y^{i-1} = a | x) = \frac{a_a(i-1)P(y^i = b | y^{i-1} = a)P(x^i | y^i = b)\beta_b(i)}{\sum_{a', b'} a_{a'}(i-1)P(y^i = b' | y^{i-1} = a')P(x^i | y^i = b')\beta_{b'}(i)}$$

# Forward (sub-)Algorithm

- Solve for every:  $\alpha_z(i) = P(x^{1:i}, y^i = Z | A, O)$

- Naively: **Exponential Time!**

$$\alpha_z(i) = P(x^{1:i}, y^i = Z | A, O) = \sum_{y^{1:i-1}} P(x^{1:i}, y^i = Z, y^{1:i-1} | A, O)$$

- Can be computed recursively (like Viterbi)

$$\alpha_z(1) = P(y^1 = z | y^0) P(x^1 | y^1 = z) = O_{x^1, z} A_{z, start}$$

$$\alpha_z(i+1) = O_{x^{i+1}, z} \sum_{j=1}^L \alpha_j(i) A_{z,j}$$

Viterbi effectively replaces sum with max

# Backward (sub-)Algorithm

- Solve for every:  $\beta_z(i) = P(x^{i+1:M} \mid y^i = Z, A, O)$

- Naively: **Exponential Time!**

$$\beta_z(i) = P(x^{i+1:M} \mid y^i = Z, A, O) = \sum_{y^{i+1:L}} P(x^{i+1:M}, y^{i+1:L} \mid y^i = Z, A, O)$$

- Can be computed recursively (like Viterbi)

$$\beta_z(M) = 1$$

$$\beta_z(i) = \sum_{j=1}^L \beta_j(i+1) A_{j,z} O_{x^{i+1}, j}$$

# Forward-Backward Algorithm

- Runs Forward       $\alpha_z(i) = P(x^{1:i}, y^i = Z | A, O)$
- Runs Backward       $\beta_z(i) = P(x^{i+1:M} | y^i = Z, A, O)$
- For each training  $x = (x^1, \dots, x^M)$ 
  - Computes each  $P(y^i)$  for  $y = (y^1, \dots, y^M)$

$$P(y^i = z | x) = \frac{a_z(i)\beta_z(i)}{\sum_{z'} a_{z'}(i)\beta_{z'}(i)}$$

# Recap: Unsupervised Training

- Train using only word sequences:  $S = \{x_i\}_{i=1}^N$   


Word Sequence  
(Sentence)
- y's are “hidden states”
  - All pairwise transitions are through y's
  - Hence hidden Markov Model
- Train using EM algorithm
  - Converge to local optimum

# Initialization

- How to choose #hidden states?
  - By hand
  - Cross Validation
    - $P(x)$  on validation data
    - Can compute  $P(x)$  via forward algorithm:

$$P(x) = \sum_y P(x, y) = \sum_z \alpha_z(M) P(\text{End} \mid y^M = z)$$

# Recap: Sequence Prediction & HMMs

- Models pairwise dependences in sequences

$x = \text{"I fish often"}$

**POS Tags:**

Det, Noun, Verb, Adj, Adv, Prep

**Independent:** (N, N, Adv)

**HMM Viterbi:** (N, V, Adv)

- Compact: only model pairwise between  $y$ 's
- **Main Limitation:** Lots of independence assumptions
  - Poor predictive accuracy

# Next Week

- Tuesday: Hidden Markov Models
  - (Unstructured Lecture)
- Thursday: Deep Generative Models
  - Recent Applications
- Recitation **Next TUESDAY (7pm):**
  - Recap of Viterbi and Forward/Backward